

**HUGGING FACE ACCELERATE: MAKING  
DEVICE-AGNOSTIC ML TRAINING AND  
INFERENCE EASY AT SCALE**

# WHO AM I?

- Zachary Mueller
- Technical Lead for the 🙌 Accelerate project
- Maintain the `transformers` Trainer
- API design geek

# WHAT IS 🤗 ACCELERATE?

- A training framework
- An inference framework
- A command-line interface

# A TRAINING FRAMEWORK

- Powered by PyTorch
- Change a few lines of code, gain device *and* hardware-agnostic capabilities
- Low-code, with minimal magic aimed at easy hackability and use without high-level abstractions
- We handle the intricacies so you don't have to

# A TRAINING FRAMEWORK

- Support for any hardware-accelerator on the market:
  - CPU, GPU, TPU, XPU, NPU, MLU
- Automatic mixed-precision training *safely* in whatever fashion you may choose:
  - FP16, BF16, FP8 (through either `TransformerEngine` or `MS-AMP`)
- Automatic and efficient gradient accumulation
- Support for quantization through `bitsandbytes`
- Support your favorite experiment trackers (`aim`, `clearml`, `comet_ml`, `dvc-lite`, `ml-flow`, `tensorboard`, `wandb`)
- Easy to configure plugin or YAML-level API for setting up advanced frameworks like `FSDP`, `DeepSpeed`, and `Megatron-LM`

# LOW-CODE

- Biggest friction with “wrapper” libraries is control of your code
- By being minimally intrusive, your code just “works” while still giving you complete control

```
1  import torch
2  import torch.nn.functional as F
3  from datasets import load_dataset
4  + from accelerate import Accelerator
5
6  + accelerator = Accelerator()
7  - device = 'cpu'
8  + device = accelerator.device
9
10 model = torch.nn.Transformer().to(device)
11 optimizer = torch.optim.Adam(model.parameters())
12 dataset = load_dataset('my_dataset')
13 data = torch.utils.data.DataLoader(dataset, shuffle=True)
14
15 + model, optimizer, dataloader = accelerator.prepare(model, optimizer, dataloader)
16
17 model.train()
18 for epoch in range(10):
19     for source, targets in dataloader:
20         source, targets = source.to(device), targets.to(device)
21         optimizer.zero_grad()
22         output = model(source)
23         loss = F.cross_entropy(output, targets)
24 -         loss.backward()
25 +         accelerator.backward(loss)
26         optimizer.step()
```

# EASY TO INTEGRATE

- Due to the low-code nature, it's trivial to integrate into existing PyTorch frameworks:

## 1. Create an `Accelerator`

```
1  import torch
2  import torch.nn.functional as F
3  from datasets import load_dataset
4  + from accelerate import Accelerator
5
6  + accelerator = Accelerator()
7  device = 'cpu'
8
9  model = torch.nn.Transformer().to(device)
10 optimizer = torch.optim.Adam(model.parameters())
11 dataset = load_dataset('my_dataset')
12 data = torch.utils.data.DataLoader(dataset, shuffle=True)
13
14 model.train()
15 for epoch in range(10):
16     for source, targets in dataloader:
17         source, targets = source.to(device), targets.to(device)
18         optimizer.zero_grad()
19         output = model(source)
20         loss = F.cross_entropy(output, targets)
21         loss.backward()
22         optimizer.step()
```

# EASY TO INTEGRATE

- Due to the low-code nature, it's trivial to integrate into existing PyTorch frameworks:

2. Wrap your PyTorch objects with `accelerator.prepare` and remove device-placements

```
1  import torch
2  import torch.nn.functional as F
3  from datasets import load_dataset
4  from accelerate import Accelerator
5
6  accelerator = Accelerator()
7  - device = 'cpu'
8
9  model = torch.nn.Transformer().to(device)
10 optimizer = torch.optim.Adam(model.parameters())
11 dataset = load_dataset('my_dataset')
12 data = torch.utils.data.DataLoader(dataset, shuffle=True)
13
14 + model, optimizer, dataloader = accelerator.prepare(model, optimizer, dataloader)
15
16 model.train()
17 for epoch in range(10):
18     for source, targets in dataloader:
19         source, targets = source.to(device), targets.to(device)
20         optimizer.zero_grad()
21         output = model(source)
22         loss = F.cross_entropy(output, targets)
23         loss.backward()
24         optimizer.step()
```



# EASY TO INTEGRATE

- Due to the low-code nature, it's trivial to integrate into existing PyTorch frameworks:

## 3. Use `accelerator.backward` for the backward pass

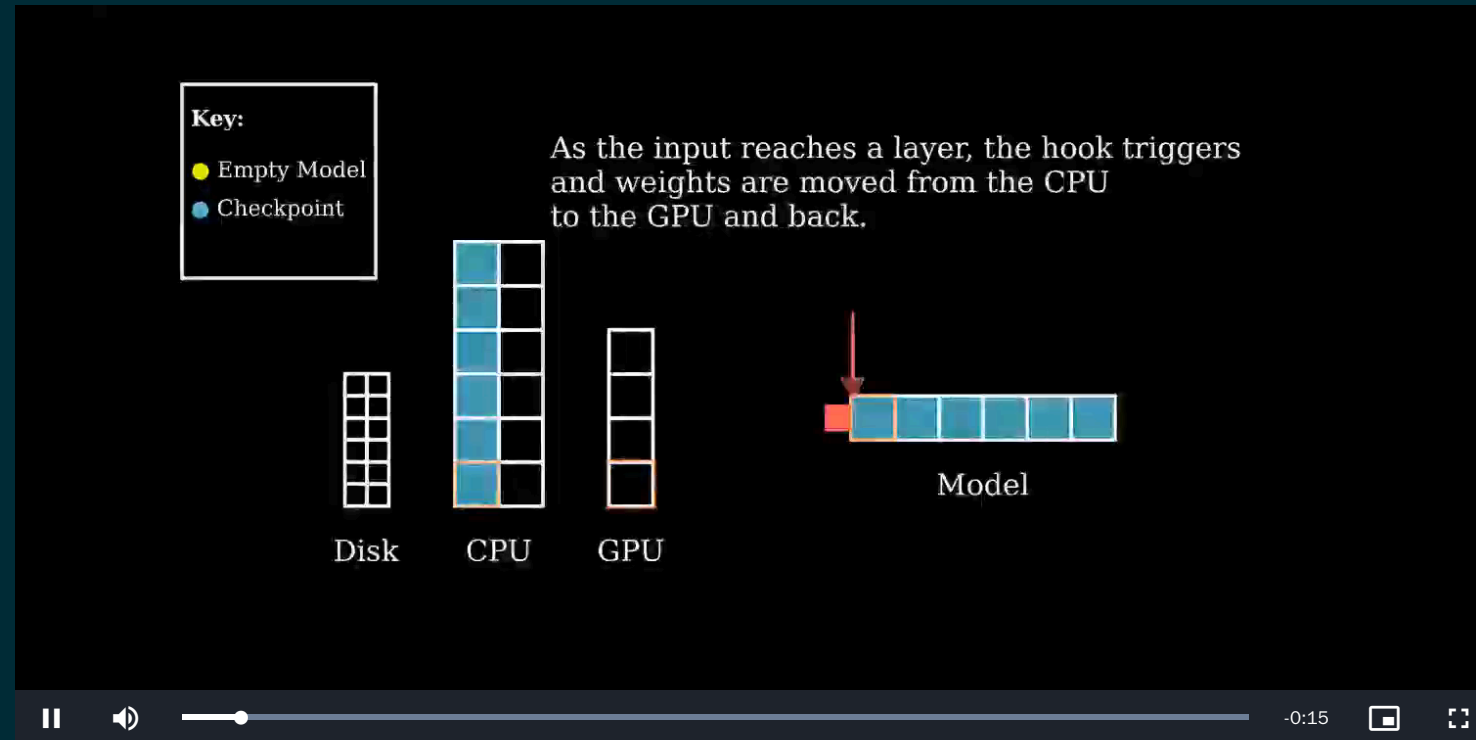
```
1  import torch
2  import torch.nn.functional as F
3  from datasets import load_dataset
4  from accelerate import Accelerator
5
6  accelerator = Accelerator()
7
8  model = torch.nn.Transformer().to(device)
9  optimizer = torch.optim.Adam(model.parameters())
10 dataset = load_dataset('my_dataset')
11 data = torch.utils.data.DataLoader(dataset, shuffle=True)
12
13 model, optimizer, dataloader = accelerator.prepare(model, optimizer, dataloader)
14
15 model.train()
16 for epoch in range(10):
17     for source, targets in dataloader:
18         source, targets = source.to(device), targets.to(device)
19         optimizer.zero_grad()
20         output = model(source)
21         loss = F.cross_entropy(output, targets)
22 -         loss.backward()
23 +         accelerator.backward(loss)
24         optimizer.step()
```

# BUT WHAT ABOUT INFERENCE?

- 🤗 Accelerate is not just for training, and has helped make the GPU-Poor take control of the narrative
- Using tools like Big Model Inference, users with *tiny* compute can run large models locally
- Started with the boom of stable diffusion, and now has scaled to having the ability to run huge LLMs locally with a single graphics card

# HOW DOES IT WORK?

- PyTorch introduced `device="meta"`
- 🙌 Accelerate introduced `device_map="auto"`



# A CLI INTERFACE

- `accelerate config`
  - Configure the environment
- `accelerate launch`
  - How to run your script

# LAUNCHING DISTRIBUTED TRAINING IS HARD

```
1 python script.py
```

VS.

```
1 torchrun --nnodes=1 --nproc_per_node=2 script.py
```

VS.

```
1 deepspeed --num_gpus=2 script.py
```

How can we make this better?

# accelerate launch

```
1 accelerate launch script.py
```

```
1 accelerate launch --multi_gpu --num_processes 2 script.py
```

```
1 accelerate launch \  
2   --multi_gpu \  
3   --use_deepspeed \  
4   --num_processes 2 \  
5   script.py
```

# accelerate config

- Rely on `config.yaml` files
- Choose to either running `accelerate config` or write your own:

`ddp_config.yaml`

```
1 compute_environment: LOCAL_MACHINE
2 distributed_type: MULTI_GPU
3 main_training_function: main
4 mixed_precision: bf16
5 num_machines: 1
6 num_processes: 8
```

`fsdp_config.yaml`

```
1 compute_environment: LOCAL_MACHINE
2 distributed_type: FSDP
3 fsdp_config:
4   fsdp_auto_wrap_policy: TRANSFORMER_BASED_WRAP
5   fsdp_backward_prefetch: BACKWARD_PRE
6   fsdp_cpu_ram_efficient_loading: true
7   fsdp_forward_prefetch: false
8   fsdp_offload_params: false
9   fsdp_sharding_strategy: FULL_SHARD
10  fsdp_state_dict_type: SHARDED_STATE_DICT
11  fsdp_sync_module_states: true
12  fsdp_use_orig_params: false
13 main_training_function: main
14 mixed_precision: bf16
15 num_machines: 1
16 num_processes: 8
```

**NOW THAT YOU'RE UP TO SPEED,  
WHAT'S NEW?**



**WE'VE HAD A BUSY LAST YEAR, AND SO  
HAS THE ML COMMUNITY!**

# NEW TRAINING TECHNIQUES

- Quantization has taken the field by storm
- New ideas such as FSDP + QLoRA to train huge models on tiny compute!
- New precision backends as we train natively on smaller precision
- Optimizing further how much we can push on a single machine through efficient RAM and timing techniques

# LARGER COMPUTE LANDSCAPE

- As we search for alternatives to NVIDIA, new compilers rise:
  - XPU (Intel)
  - NPU (Intel)
  - MLU (Cambricon)

All of which are supported by 🤖 Accelerate

# LOWER ABSTRACTIONS

- While the `Accelerator` was great, needed better abstractions focused on controlling behaviors
- Introduced the `PartialState`

```
1  from accelerate import PartialState
2
3  if PartialState().is_main_process:
4      # Run on only 1 device
5
6  with PartialState().main_process_first:
7      # Useful for dataset processing
8
9  # Device-agnostic without the bulk of the `Accelerator`
10 device = PartialState().device
```

# FASTER AND BETTER INFERENCE ALTERNATIVES

- PiPPy gives us efficient pipeline-parallelism in distributed environments to increase throughput while keeping a simple torch-bound API
- Rather than having to wait for each GPU, every GPU can be busy in parallel

```
1 import torch
2 from transformers import AutoModelForSequenceClassification
3
4 from accelerate import PartialState, prepare_pippy
5
6 model = AutoModelForSequenceClassification.from_pretrained("gpt2")
7 model.eval()
8
9 input = torch.randint(
10     low=0,
11     high=model.config.vocab_size,
12     size=(2, 1024), # bs x seq_len
13     device="cpu",
14 )
15
16 model = prepare_pippy(model, split_points="auto", example_args=(input,
17
18 with torch.no_grad():
19     output = model(input)
```


# **ADOPTION: ACCELERATE IN THE ECOSYSTEM**

# ACCELERATE IN THE ECOSYSTEM

- Many of the frameworks you use daily already rely on 🙌  
Accelerate!
  - Nearly all of 🙌
  - `axolotl`
  - `fastai`
  - FastChat
  - `lucidrains`
  - `kornia`

# ACCELERATE IN THE ECOSYSTEM

- Started as a way to isolate out distributed code on TPU and `DistributedDataParallelism`

 **Sylvain Gugger**  
@GuggerSylvain

Want to run your PyTorch training loop on multi-GPUs or TPUs without using an abstract class you can't control or tweak easily? Try out 🤖 Accelerate!

```
+ accelerator = Accelerator()
- device = 'cpu'

+ model = torch.nn.Transformer()
- model = torch.nn.Transformer().to(device)
  optim = torch.optim.Adam(model.parameters())

dataset = load_dataset('my_dataset')
data = torch.utils.data.DataLoader(dataset, shuffle=True)

+ model, optim, data = accelerator.prepare(model, optim, data)

model.train()
for epoch in range(10):
    for source, targets in data:
        - source = source.to(device)
        - targets = targets.to(device)

        optimizer.zero_grad()

        output = model(source)
        loss = F.cross_entropy(output, targets)

        + accelerator.backward(loss)
        loss.backward()
```

Introducing 🤖 Accelerate

From [huggingface.co](https://huggingface.co)

11:33 AM · Apr 16, 2021



# ACCELERATE IN THE ECOSYSTEM

- Now is the backbone of some of the largest PyTorch training frameworks in the ecosystem



**WHAT'S NEXT?**

# ELEVATING THE COMMUNITY

- Now that more advanced training techniques are reachable (FSDP, DeepSpeed, etc), we need to focus on educating the community on how to use it best
- Goes beyond how to use the `Trainer` or `Accelerator`, but how to use *what* where
- Keep Accelerate as a tool for the community to utilize when new techniques come out and play with, to push new ideas to scale quickly

# 1.0.0: SOON!

- Tried and battle-tested by over 7M users/month
- As we've been stable for over a year now, we're near ready to release 1.0.0

# THANKS FOR JOINING!

- 🙌 Accelerate documentation
- Launching distributed code
- Distributed code and Jupyter Notebooks
- Migrating to 🙌 Accelerate easily
- Big Model Inference tutorial
- DeepSpeed and 🙌 Accelerate
- Fully Sharded Data Parallelism and 🙌 Accelerate
- FSDP vs DeepSpeed In-Depth